

IST Lab 02: LOGICAL VOLUMES AND SNAPSHOTS

Authors: Kevin Auberson and Paul Gillet

Group: L1GrG **Note:** L2GrH

Date: March 7, 2024

TASK 1: CREATE PHYSICAL VOLUMES

1. Reset your external disk. Using parted remove all partitions, or simply write a new partition table.

```
$ sudo parted
GNU Parted 3.4
Using /dev/sda
Welcome to GNU Parted! Type 'help' to view a list of commands.
(parted) select /dev/sdb
Using /dev/sdb
(parted) print
Model: Generic Flash Disk (scsi)
Disk /dev/sdb: 31.5GB
Sector size (logical/physical): 512B/512B
Partition Table: msdos
Disk Flags:

Number  Start   End     Size    Type    File system  Flags
  1      1049kB  15.7GB  15.7GB  primary fat32         lba
  2      15.7GB  31.5GB  15.7GB  primary ext4

(parted) rm 1
(parted) rm 2
```

2. Create four partitions with these characteristics: primary, 25 MB size, type ext4.

```
(parted) mkpart primary ext4 0MB 25MB
(parted) mkpart primary ext4 25MB 50MB
(parted) mkpart primary ext4 50MB 75MB
```

```
(parted) mkpart primary ext4 75MB 100MB
(parted) p
Model: Generic Flash Disk (scsi)
Disk /dev/sdb: 31.5GB
Sector size (logical/physical): 512B/512B
Partition Table: msdos
Disk Flags:
```

Number	Start	End	Size	Type	File system	Flags
1	512B	25.0MB	25.0MB	primary		
2	25.2MB	50.3MB	25.2MB	primary		
3	50.3MB	75.5MB	25.2MB	primary		
4	75.5MB	99.6MB	24.1MB	primary		

3. List the available LVM commands. They belong to the Debian package lvm2 which should already be installed. Use dpkg with the -L option to list the content of the package. The commands are all located in the /sbin directory. Use grep to filter and sort to sort alphabetically.

```
$ dpkg -L lvm2 | grep /sbin | sort
/sbin
/sbin/fsadm
/sbin/lvchange
/sbin/lvconvert
/sbin/lvcreate
/sbin/lvdisplay
/sbin/lvextend
/sbin/lvm
/sbin/lvmconfig
/sbin/lvmdiskscan
/sbin/lvmdump
/sbin/lvmpolld
/sbin/lvmsadc
/sbin/lvmsar
/sbin/lvreduce
/sbin/lvremove
/sbin/lvrename
/sbin/lvresize
/sbin/lvs
/sbin/lvscan
```

```
/sbin/pvchange
/sbin/pvck
/sbin/pvcreate
/sbin/pvdisplay
/sbin/pvmove
/sbin/pvremove
/sbin/pvresize
/sbin/pvs
/sbin/pvscan
/sbin/vgcfgbackup
/sbin/vgcfgrestore
/sbin/vgchange
/sbin/vgck
/sbin/vgconvert
/sbin/vgcreate
/sbin/vgdisplay
/sbin/vgexport
/sbin/vgextend
/sbin/vgimport
/sbin/vgimportclone
/sbin/vgmerge
/sbin/vgmknodes
/sbin/vgreduce
/sbin/vgremove
/sbin/vgrename
/sbin/vgs
/sbin/vgscan
/sbin/vgsplit
```

4. List all partitions that could potentially host a Physical Volume by using pvs with the --all option.

```
$ sudo pvs --all
PV          VG Fmt Attr PSize PFree
/dev/loop1          ---    0    0
/dev/loop10         ---    0    0
/dev/loop11         ---    0    0
/dev/loop13         ---    0    0
/dev/loop2          ---    0    0
/dev/loop3          ---    0    0
```

/dev/loop4	---	0	0
/dev/loop5	---	0	0
/dev/loop6	---	0	0
/dev/loop7	---	0	0
/dev/loop8	---	0	0
/dev/loop9	---	0	0
/dev/sda2	---	0	0
/dev/sda3	---	0	0
/dev/sdb1	---	0	0
/dev/sdb2	---	0	0
/dev/sdb3	---	0	0
/dev/sdb4	---	0	0

5. On the four partitions of your external disk, create four Physical Volumes using `pvcreate`. Add the `-vv` option so that it tells you in detail what it is doing. For the first partition copy the output of the command into the report, but copy only the lines about the partition that receives the Physical Volume and ignore the other messages.

```
$ sudo pvcreate /dev/sdb1 -vv

/dev/sdb1: using cached size 48828 sectors
/dev/sdb1: using cached size 48828 sectors
/dev/sdb1: No lvm label detected
/dev/sdb1: using cached size 48828 sectors
/dev/sdb1: using cached size 48828 sectors
/dev/sdb1: No lvm label detected
Wiping signatures on new PV /dev/sdb1.
/dev/sdb1: using cached size 48828 sectors
devices/default_data_alignment not found in config: defaulting to 1
Device /dev/sdb1: queue/minimum_io_size is 512 bytes.
Device /dev/sdb1: queue/optimal_io_size is 0 bytes.
Device /dev/sdb1: alignment_offset is 0 bytes.
Set up physical volume for "/dev/sdb1" with 48828 available sectors.
Scanning for labels to wipe from /dev/sdb1
Zeroing start of device /dev/sdb1.
Writing physical volume data to disk "/dev/sdb1".
/dev/sdb1: Writing label to sector 1 with stored offset 32.
Physical volume "/dev/sdb1" successfully created.
Unlocking /run/lock/lvm/P_global
```

6. Display detailed information about the first Physical Volume using pvdisplay.

```
$ sudo pvdisplay /dev/sdb1
"/dev/sdb1" is a new physical volume of "23.84 MiB"
--- NEW Physical volume ---
PV Name                /dev/sdb1
VG Name
PV Size                23.84 MiB
Allocatable            NO
PE Size                0
Total PE               0
Free PE                0
Allocated PE           0
PV UUID                0e2h5r-x90G-fVst-N6oE-boxD-vree-tSR0hy
```

TASK 2: CREATE TWO VOLUME GROUPS

1. Create a first Volume Group lab-vg1 that contains only the first Physical Volume. Display the Physical Volume again with pvdisplay. What has changed?

```
$ sudo vgcreate lab-vg1 /dev/sdb1
Volume group "lab-vg1" successfully created

$ sudo pvdisplay /dev/sdb1
--- Physical volume ---
PV Name                /dev/sdb1
VG Name                lab-vg1
PV Size                23.84 MiB / not usable 3.84 MiB
Allocatable            yes
PE Size                4.00 MiB
Total PE               5
Free PE                5
Allocated PE           0
PV UUID                0e2h5r-x90G-fVst-N6oE-boxD-vree-tSR0hy
```

As we can see, the Volume Group was successfully created and the PV is now in the VG lab-vg1 also the PV is subdivided into 5 Physical Extends (PEs) with the same size of 4MB (default size)

2.Create a second Volume Group lab-vg2 that contains Physical Volumes 2 and 3.

```
$ sudo vgcreate lab-vg2 /dev/sdb2 /dev/sdb3
Volume group "lab-vg2" successfully created
```

3.List all Volume Groups with vgs. Then list all Physical Volumes with pvs. What do you see?

```
$ sudo vgs
VG          #PV #LV #SN Attr   VSize  VFree
lab-vg1     1   0   0 wz--n- 20.00m 20.00m
lab-vg2     2   0   0 wz--n- 40.00m 40.00m

$ sudo pvs
PV          VG          Fmt  Attr PSize  PFree
/dev/sdb1   lab-vg1   lvm2 a--  20.00m 20.00m
/dev/sdb2   lab-vg2   lvm2 a--  20.00m 20.00m
/dev/sdb3   lab-vg2   lvm2 a--  20.00m 20.00m
/dev/sdb4                   lvm2 ---  23.00m 23.00m
```

As we can see we have 2 VG containing 1 PV and 2 PV for the 2nd one. We can then see with the command pvs all our PVs with their VG. Only PV "/dev/sdb4" is not in a VG.

TASK 3: CREATE LOGICAL VOLUMES

1.On the Volume Group lab-vg1 create a Logical Volume of size 20 MB with the command lvcreate -L 20M lab-vg1.

```
$ sudo lvcreate -L 20M lab-vg1

Logical volume "lvol0" created.
```

2.Verify hat the new volume appears when you use lvs to list Logical Volumes. Also verify that it appears when you use lsblk to list the block devices. What is the name of the special file in /dev that represents the volume?

```
$ sudo lvs

LV      VG      Attr          LSize  Pool Origin Data%  Meta%  Move Log Cpy%Sync
Convert
```

```
lvol0 lab-vg1 -wi-a—— 20.00m
```

```
$ lsblk
```

NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPOINTS
loop0	7:0	0	63.9M	1	loop	/snap/core20/2182
loop1	7:1	0	4K	1	loop	/snap/bare/5
loop2	7:2	0	63.4M	1	loop	/snap/core20/1974
loop3	7:3	0	266.6M	1	loop	/snap/firefox/3836
loop4	7:4	0	73.9M	1	loop	/snap/core22/858
loop5	7:5	0	237.2M	1	loop	/snap/firefox/2987
loop6	7:6	0	349.7M	1	loop	/snap/gnome-3-38-2004/143
loop7	7:7	0	485.5M	1	loop	/snap/gnome-42-2204/120
loop8	7:8	0	12.3M	1	loop	/snap/snap-store/959
loop9	7:9	0	53.3M	1	loop	/snap/snapd/19457
loop10	7:10	0	40.4M	1	loop	/snap/snapd/20671
loop11	7:11	0	497M	1	loop	/snap/gnome-42-2204/141
loop12	7:12	0	91.7M	1	loop	/snap/gtk-common-themes/1535
loop13	7:13	0	452K	1	loop	/snap/snapd-desktop-integration/83
sda	8:0	0	25G	0	disk	
└sda1	8:1	0	1M	0	part	
└sda2	8:2	0	513M	0	part	/boot/efi
└sda3	8:3	0	24.5G	0	part	/var/snap/firefox/common/host-hunspell /
sdb	8:16	1	29.3G	0	disk	
└sdb1	8:17	1	23.8M	0	part	
└lab--vg1-lvol0						
	252:0	0	20M	0	lvm	
└sdb2	8:18	1	24M	0	part	
└sdb3	8:19	1	24M	0	part	
└sdb4	8:20	1	23M	0	part	
sr0	11:0	1	1024M	0	rom	

The special file is stored in `/dev/lab-vg1/lvol0`. The directory `lab-vg1` has been created when we created the logical volume in this group.

3. Create an ext4 file system on the volume. Mount the volume. Fill the file system with a 14 MB file using `dd` (Google it).

```
$ sudo mkfs.ext4 /dev/lab-vg1/lvol0

mke2fs 1.46.5 (30-Dec-2021)

Creating filesystem with 5120 4k blocks and 5120 inodes

Allocating group tables: done

Writing inode tables: done

Creating journal (1024 blocks): done

Writing superblocks and filesystem accounting information: done

$ sudo dd if=/dev/lab-vg1/lvol0 of=test bs=1M count=14
14+0 records in
14+0 records out
14680064 bytes (15 MB, 14 MiB) copied, 8.58476 s, 1.7 MB/s
```

4. On the Volume Group lab-vg2 create another Logical Volume of size 20 MB, create an ext4 file system on it and mount it. Create a file named foo that contains the text 111.

```
$ sudo lvcreate -L 20M lab-vg2

Logical volume "lvol0" created.

$ sudo mkdir /mnt/vg2-lvol0
$ sudo mkfs.ext4 /dev/lab-vg2/lvol0
mke2fs 1.46.5 (30-Dec-2021)

Creating filesystem with 5120 4k blocks and 5120 inodes

Allocating group tables: done

Writing inode tables: done

Creating journal (1024 blocks): done

Writing superblocks and filesystem accounting information: done
```



```
$ sudo mount /dev/lab-vg2/lvol0 /mnt/vg2-lvol0/
```

```
$ sudo sh -c 'echo "111" > foo'
```

TASK 4: GROW A FILE SYSTEM WHILE IT IS IN USE

1. Verify that the file system is indeed full (use `df -h`).

```
$ df -h
Filesystem                                Size  Used Avail Use% Mounted on
...
/dev/mapper/lab--vg1-lvol0                15M   15M    0 100% /mnt/lvol0
/dev/mapper/lab--vg2-lvol0                15M   32K   14M   1% /mnt/vg2-lvol0
```

2. Verify that the Volume Group is full (use `vgs`).

```
$ sudo vgs
VG      #PV #LV #SN Attr   VSize  VFree
lab-vg1    1  1  0 wz--n- 20.00m    0
lab-vg2    2  1  0 wz--n- 40.00m 20.00m
```

3. Extend the Volume group using `vgextend` and verify with `vgs`.

```
$ sudo vgextend lab-vg1 /dev/sdb4
Volume group "lab-vg1" successfully extended
$ sudo vgs
VG      #PV #LV #SN Attr   VSize  VFree
lab-vg1    2  1  0 wz--n- 40.00m 20.00m
lab-vg2    2  1  0 wz--n- 40.00m 20.00m
```

4. Extend the Logical Volume by an additional 20 MB using `lvextend --size <new_size> <volume_group>/<logical_volume>`.

4.4 La taille n'est pas étendue ici (de 20 à 20).
Il faut passer 40 pour 20 Mb supplémentaire (soit 40, soit +20). (0/1)

```
$ sudo lvextend --size 20M lab-vg1/lvol0
New size (5 extents) matches existing size (5 extents).
```

5. Grow the file system while it is mounted using `resize2fs` and verify its new capacity with `df -h`

```
$ sudo resize2fs /dev/lab-vg1/lvol0
```

```
resize2fs 1.46.5 (30-Dec-2021)
```

```
The filesystem is already 5120 (4k) blocks long.  Nothing to do!
```

```
$ df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
...					
/dev/mapper/lab--vg2-lvol0	15M	32K	14M	1%	/mnt/vg2-lvol0
/dev/mapper/lab--vg1-lvol0	15M	15M	0	100%	/mnt/lvol0

Après la commande, il est indiqué que le système de fichier est déjà à sa taille maximum, donc rien de plus n'est fait.

Oui, parce qu'elle n'a pas été étendue (cf au-dessus)!

The command output says that the filesystem is already up to it's maximum size, so there's nothing more to be done.

TASK 5: CREATE A SNAPSHOT

1. Create a snapshot volume using the `--snapshot` option of `lvcreate`. Use the `--name` option to give it the name `snap`. You also need to specify a size for the snapshot with the `--size` option. Remember that initially a snapshot does not consume any storage blocks as the data in the original volume and the snapshot volume is identical. It is only when the data in the two volumes starts deviating that storage blocks are needed. The size of the snapshot determines how many data blocks can be different

```
$ sudo lvcreate --snapshot --name snap --size 15M /dev/lab-vg2/lvol0
Rounding up size to full physical extent 16.00 MiB
Logical volume "snap" created.
```

2. Display an overview of all Logical Volumes using `lvs`. Which column shows the name of the original volume?

```
$ sudo lvs
  LV      VG      Attr      LSize   Pool Origin Data%  Meta%  Move Log Cpy%Sync
Convert
  lvol0 lab-vg1 -wi-a----- 20.00m
  lvol0 lab-vg2 owi-aos--- 20.00m
  snap lab-vg2 swi-a-s--- 16.00m          lvol0  0.07
```

The sixth column "Origin" show us the name of the original volume.

3.Display the characteristics of the snapshot volume using `lvdisplay`.

```
$ sudo lvdisplay

...

--- Logical volume ---
LV Path                /dev/lab-vg2/snap
LV Name                 snap
VG Name                 lab-vg2
LV UUID                 HPtKNz-6vZz-dedq-Bwya-el44-TpbX-ih3NNn
LV Write Access         read/write
LV Creation host, time kevin-VirtualBox, 2024-03-17 19:37:42 +0100
LV snapshot status      active destination for lvol0
LV Status                available
# open                  0
LV Size                  20.00 MiB
Current LE               5
COW-table size           16.00 MiB
COW-table LE             4
Allocated to snapshot    0.07%
Snapshot chunk size      4.00 KiB
Segments                 1
Allocation                inherit
Read ahead sectors       auto
- currently set to       256
Block device             252:4

...
```

3.1.What line shows the name of the original volume?

The "LV Path" shows the name of the original volume

3.2.What line shows the size of the original volume?

The ninth lines "LV Size" shows the size of the original volume.

3.3.What line shows the space allocated for the snapshot volume?

3.4.What does COW stand for?

COW means "Copy-On-Write", it is used when a snapshot is taken. Instead of copying all data, it copies only the metadata like pointers to data blocks. This means that the actual blocks are shared between the original volume and the snapshot. Whenever a block is modified in the original volume, the original data block is first copied to the snapshot, and then the modification is made. This allows to write only when changes are made to new blocks, the original data are preserved on both the snapshot and original volume.

4. Mount the snapshot volume. Using the file foo you created earlier verify that the two volumes behave like independent copies.

```
$ sudo mkdir snap
$ sudo mount /dev/lab-vg2/snap /mnt/snap
$ sudo sh -c 'echo "test" > foo'
$ cat foo
test
$ cat /mnt/vg2/foo
111
```

As we can see changes are not yet written

5.Make the data of the original volume change completely by using the dd command to write a new file of 14 MB size. Run df -h to see how it affects the fullness of the original volume and the snapshot. What do you see?

```
$ sudo dd if=/dev/lab-vg2/lvol0 of=newtest bs=1M count=14
14+0 records in
14+0 records out
14680064 bytes (15 MB, 14 MiB) copied, 14.4286 s, 1.0 MB/s

$ df -h
Filesystem                Size      Used Avail Use% Mounted on
tmpfs                      794M    1.7M   792M   1% /run
```

/dev/sda3	24G	13G	11G	55%	/
tmpfs	3.9G	0	3.9G	0%	/dev/shm
tmpfs	5.0M	4.0K	5.0M	1%	/run/lock
/dev/sda2	512M	6.1M	506M	2%	/boot/efi
tmpfs	794M	92K	794M	1%	/run/user/1000
/dev/mapper/lab--vg2-lvol0	15M	15M	0	100%	/mnt/vg2
/dev/mapper/lab--vg2-snap	15M	28K	14M	1%	/mnt/snap

As we can see the original volume is full, however the snap volume is unaffected.

5.1.The way that you allocated it, is the snapshot volume able support a change of 14 MB of data?

The snapshot is able to support change of 14MB of data because it only stores the changes made since the snapshot was taken.

5.2.What happened? Why?

The snapshot volume remained unchanged because it only tracks the changes made to the original volume, rather than storing duplicate data.

6.Remove the broken snapshot volume.

```
$ sudo lvremove /dev/lab-vg2/snap
Do you really want to remove and DISCARD active logical volume lab-vg2/snap?
[y/n]: y
Logical volume "snap" successfully removed
```

7.Redo the above, this time allocating sufficient space to the snapshot volume to support a complete change of data of the original volume.

```
$ sudo lvcreate --snapshot --name snap --size 20M /dev/lab-vg2/lvol0
Logical volume "snap" created.

$ sudo lvs
LV      VG      Attr      LSize   Pool Origin Data%  Meta%  Move Log Cpy%Sync
Convert
lvol0   lab-vg1 -wi-ao--- 20.00m
lvol0   lab-vg2 owi-aos--- 20.00m
snap    lab-vg2 swi-a-s--- 20.00m          lvol0 0.06
```

```
$ sudo lvdisplay
```

```
...
```

```
--- Logical volume ---
```

```
LV Path                /dev/lab-vg2/snap
LV Name                 snap
VG Name                 lab-vg2
LV UUID                 XiiyxJ-dnH9-0fXz-KOVU-Uu4W-Gp01-Q4FgzD
LV Write Access         read/write
LV Creation host, time kevin-VirtualBox, 2024-03-21 09:41:20 +0100
LV snapshot status      active destination for lvol0
LV Status                available
# open                  0
LV Size                  20.00 MiB
Current LE               5
COW-table size           20.00 MiB
COW-table LE             5
Allocated to snapshot    0.06%
Snapshot chunk size      4.00 KiB
Segments                 1
Allocation               inherit
Read ahead sectors       auto
- currently set to      256
Block device             252:4
```

```
...
```

```
$ sudo mkdir /mnt/snap
```

```
$ sudo mount /dev/lab-vg2/snap /mnt/snap
```

```
$ sudo sh -c 'echo "test" > foo'
```

```
$ cat foo
```

```
111
```

```
$ cat /mnt/vg2-lvol0/foo
```

```
test
```

```
$ sudo dd if=/dev/lab-vg2/lvol0 of=newtest bs=1M count=14
```

```
14+0 records in
```

```
14+0 records out
```

```
14680064 bytes (15 MB, 14 MiB) copied, 15.7727 s, 931 kB/s
```

```
$ df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
tmpfs	196M	1.5M	195M	1%	/run
/dev/sda3	24G	12G	11G	53%	/
tmpfs	980M	0	980M	0%	/dev/shm
tmpfs	5.0M	4.0K	5.0M	1%	/run/lock
/dev/sda2	512M	6.1M	506M	2%	/boot/efi
tmpfs	196M	92K	196M	1%	/run/user/1000
/dev/mapper/lab--vg2-lvol0	15M	15M	0	100%	/mnt/vg2-lvol0
/dev/mapper/lab--vg2-snap	15M	28K	14M	1%	/mnt/snap

TASK 6: PROVISION A THIN VOLUME AND SNAPSHOT IT

1.Remove all Logical Volumes from Volume Group lab-vg2.

```
$ sudo umount /dev/lab-vg2/lvol0

$ sudo lvchange -an /dev/lab-vg2/*

$ sudo lvremove lab-vg2
Do you really want to remove and DISCARD logical volume lab-vg2/lvol0? [y/n]:
y
Logical volume "lvol0" successfully removed
```

2.Follow the explanations in the Ubuntu manual on lvmthin to create

- 2.1.a thin data Logical Volume called pool0 of 28 MB
- 2.2.a thin metadata Logical Volume called pool0meta of 4 MB

```
$ sudo lvcreate -n pool0 -L 28M lab-vg2
WARNING: ext4 signature detected on /dev/lab-vg2/pool0 at offset 1080. Wipe
it? [y/n]: y
Wiping ext4 signature on /dev/lab-vg2/pool0.
Logical volume "pool0" created.

$ sudo lvcreate -n pool0meta -L 4M lab-vg2
Logical volume "pool0meta" created.
```

3.Combine the two into a thin pool Logical Volume. List the Logical Volumes using lvs. Use the -a option to list also the hidden ones.

```
$ sudo lvconvert --type thin-pool --poolmetadata lab-vg2/pool0meta lab-vg2/pool0
```

Thin pool volume with chunk size 64.00 KiB can address at most 15.81 TiB of data.

WARNING: Converting lab-vg2/pool0 and lab-vg2/pool0meta to thin pool's data and metadata volumes with metadata wiping.

THIS WILL DESTROY CONTENT OF LOGICAL VOLUME (filesystem etc.)

Do you really want to convert lab-vg2/pool0 and lab-vg2/pool0meta? [y/n]: y

Converted lab-vg2/pool0 and lab-vg2/pool0meta to thin pool.

```
$ sudo lvs -a
```

LV	VG	Attr	LSize	Pool	Origin	Data%	Meta%	Move	Log
lvol0	lab-vg1	-wi-a-----	20.00m						
[lvol0_pmspare]	lab-vg2	ewi-----	4.00m						
pool0	lab-vg2	twi-a-tz--	28.00m			0.00	10.84		
[pool0_tdata]	lab-vg2	Twia-ao---	28.00m						
[pool0_tmeta]	lab-vg2	ewia-ao---	4.00m						

4. Create a thin Logical Volume from the thin pool named thin1 and give it a size of 80 MB, although the thin pool only has 28 MB capacity. What warnings do you see?

```
$ sudo lvcreate -n thin1 -V 80M --thinpool pool0 lab-vg2
```

WARNING: Sum of all thin volume sizes (80.00 MiB) exceeds the size of thin pool lab-vg2/pool0 and the size of whole volume group (40.00 MiB).

WARNING: You have not turned on protection against thin pools running out of space.

WARNING: Set activation/thin_pool_autoextend_threshold below 100 to trigger automatic extension of thin pools before they get full.

Logical volume "thin1" created.

The following warnings mean:

The first one: indicates that the total size of thin volumes requested (80.00 MiB) exceeds both the size of the thin pool ("pool0") and the size of the entire volume group ("lab-vg2")

The second one : means that there is no protection to prevent thin pools from running out of space. It's then possible for it to become full, it might lead to data loss or disruption of operations that depend on the thin pool

The third one : advises setting the activation/thin_pool_autoextend_threshold parameter below 100 to trigger automatic extension of thin pools before they become full.

5. Create an ext4 file system on thin1. Mount the file system. How much capacity does df -h see in the file system?

```
$ sudo mkfs.ext4 /dev/lab-vg2/thin1
mke2fs 1.46.5 (30-Dec-2021)
Discarding device blocks: done
Creating filesystem with 20480 4k blocks and 20480 inodes

Allocating group tables: done
Writing inode tables: done
Creating journal (1024 blocks): done
Writing superblocks and filesystem accounting information: done

$ df -h
```

Filesystem	Size	Used	Avail	Use%	Mounted on
tmpfs	196M	1.5M	195M	1%	/run
/dev/sda3	24G	12G	11G	53%	/
tmpfs	980M	0	980M	0%	/dev/shm
tmpfs	5.0M	4.0K	5.0M	1%	/run/lock
/dev/sda2	512M	6.1M	506M	2%	/boot/efi
tmpfs	196M	92K	196M	1%	/run/user/1000
/dev/mapper/lab--vg2--thin1	71M	24K	66M	1%	/mnt/thin

On df -h see 71M capacity for this file system.

6. Do experiments: Fill the file system with a bit of data by using dd to write files and verify that it behaves normally. Then write more and more data until you cross the size of the thin pool and see what happens. You can see LVM's log messages by using the dmesg command, they appear as device-mapper. What do you observe?

```
$ sudo dd if=/dev/zero of=/mnt/lab-vg2/thin1/testfile bs=1M count=5
5+0 records in
5+0 records out
5242880 bytes (5.2 MB, 5.0 MiB) copied, 0.0058668 s, 894 MB/s
```

```
$ sudo dd if=/dev/zero of=/mnt/lab-vg2/thin1/testfile2 bs=1M count=10
10+0 records in
10+0 records out
10485760 bytes (10 MB, 10 MiB) copied, 0.0120071 s, 873 MB/s

$ sudo dmesg | grep device-mapper

[ 0.527830] device-mapper: core: CONFIG_IMA_DISABLE_HTABLE is disabled.
Duplicate IMA measurements will not be recorded in the IMA log.

[ 0.527839] device-mapper: uevent: version 1.0.3

[ 0.528026] device-mapper: ioctl: 4.48.0-ioclt (2023-03-01) initialised:
dm-devel@redhat.com

[ 2793.954229] device-mapper: thin: Data device (dm-2) discard unsupported:
Disabling discard passdown.

[ 2826.556547] device-mapper: thin: Data device (dm-2) discard unsupported:
Disabling discard passdown.

[ 4334.735831] device-mapper: thin: Data device (dm-2) discard unsupported:
Disabling discard passdown.

[ 4335.085779] device-mapper: thin: Data device (dm-2) discard unsupported:
Disabling discard passdown.

[ 5121.494862] device-mapper: thin: 252:3: reached low water mark for data
device: sending event.

[ 5121.516711] device-mapper: thin: 252:3: switching pool to out-of-data-space
(queue IO) mode

[ 5182.426473] device-mapper: thin: 252:3: switching pool to out-of-data-space
(error IO) mode
```

device-mapper: thin: Data device (dm-2) discard unsupported: Disabling discard passdown :
This message indicates that the thin provisioned pool doesn't support the discard operation,
which is a mechanism for releasing unused data blocks.

device-mapper: thin: reached low water mark for data device: sending event : This message indicate that available space is running low.

device-mapper: thin: switching pool to out-of-data-space (queue IO) mode: This message indicates that the thin provisioned pool is switching to out-of-data-space mode, which implies that it's queuing IO operations.

These messages gives us informations on how the pool initially handles writes normally, but as it nears capacity, it starts queuing IO operations and eventually enters an error state, resulting in failed IO operations.

TASK 7: SCENARIO

You are a data engineer in a company. You need to set up the backup system for your production system, which runs a large database. Backups are important in case of unexpected incidents or human error. As the volume of data is significant, the backup process's duration can be fairly long, typically between 30 minutes and 1 hour.

Design a solution for this backup system using snapshots, knowing that:

- The backup should be performed without interrupting the database operations.
- The backup files need to be physically distant and sent to another datacenter located a few kilometers away.

Describe a potential solution and explain your thought process

A simple solution we can implement is by using LVM like we did in this lab to create snapshots of the DB. Those snapshots shall be scheduled to be made during off peak hours to minimize impact on performance of the DB.

Then transfert over SSH those snapshots to another datacenter located kilometers away to ensure data integrity during the transit.

Using monitoring tools we can ensure that the creation process is done successfully and track data transfer.

Using the 3-2-1 rule that is a data protection strategy that recommends having three copies of your data, stored on two different types of media, with one copy kept off-site.

We can also install an appliance of NetBackup to manage the backup.

(4/6): Good! Missing aspect: Initial full backup, delta changes and incremental backup